

# Deep Q-Networks (DQN): A Detailed Overview

Your Name

July 3, 2024

## 1 Introduction

Deep Q-Networks (DQN) is a reinforcement learning algorithm that combines Q-learning with deep neural networks. It was introduced by Mnih et al. in 2015 and has been successfully applied to a variety of tasks, including playing Atari games from raw pixel inputs.

## 2 DQN Algorithm

The main idea of DQN is to approximate the Q-function, which represents the expected cumulative reward for taking an action  $a$  in a state  $s$  and following the optimal policy thereafter, using a neural network. The Q-function is updated using the Bellman equation:

$$Q(s, a) = r + \gamma \max_{a'} Q(s', a')$$

where:

- $Q(s, a)$  is the estimated Q-value for taking action  $a$  in state  $s$ .
- $r$  is the reward received after taking action  $a$  in state  $s$ .
- $\gamma$  is the discount factor, which determines the importance of future rewards.
- $s'$  is the next state.

## 3 Replay Buffer

The replay buffer is a crucial component of DQN that stores the agent's experiences, allowing the algorithm to break the correlations between consecutive samples and improve the efficiency and stability of the learning process.

### 3.1 How Replay Buffer Works

1. **Storage:** The replay buffer stores tuples of the form  $(s, a, r, s')$  (state, action, reward, next state).
2. **Sampling:** During training, a random mini-batch of experiences is sampled from the replay buffer to update the Q-network. This random sampling breaks the correlation between consecutive experiences.
3. **Capacity:** The replay buffer has a fixed capacity. When the buffer is full, the oldest experiences are discarded to make room for new ones.

## 4 Parameters and Their Utility

- **Learning Rate ( $\alpha$ ):** Determines the step size for updating the Q-network weights. A higher learning rate can speed up learning but may cause instability.
- **Discount Factor ( $\gamma$ ):** Determines the importance of future rewards. A value close to 1 makes the agent consider long-term rewards, while a value close to 0 makes it focus on immediate rewards.
- **Batch Size:** The number of experiences sampled from the replay buffer for each training step. Larger batch sizes provide more stable updates but require more memory and computation.
- **Replay Buffer Size:** The maximum number of experiences the replay buffer can hold. A larger buffer can provide more diverse experiences but requires more memory.
- **Exploration Rate ( $\epsilon$ ):** Controls the trade-off between exploration (choosing random actions) and exploitation (choosing the best-known action). The exploration rate typically decays over time.
- **Target Network Update Frequency:** Determines how often the target network is updated with the weights of the Q-network. Less frequent updates stabilize training but slow down learning.

### 4.1 Algorithm

1. Initialize replay buffer  $D$ .
2. Initialize Q-network  $Q(s, a; \theta)$  with random weights.
3. Initialize target network  $Q'(s, a; \theta^-) = Q(s, a; \theta)$ .
4. For each episode:
  - (a) Initialize state  $s$ .

- (b) For each step in the episode:
- i. With probability  $\epsilon$ , select a random action  $a$ , otherwise select  $a = \arg \max_a Q(s, a; \theta)$ .
  - ii. Execute action  $a$  and observe reward  $r$  and next state  $s'$ .
  - iii. Store transition  $(s, a, r, s')$  in replay buffer  $D$ .
  - iv. Sample random mini-batch of transitions  $(s_j, a_j, r_j, s'_j)$  from  $D$ .
  - v. Set target for each mini-batch transition:

$$y_j = \begin{cases} r_j & \text{if } s'_j \text{ is terminal} \\ r_j + \gamma \max_{a'} Q(s'_j, a'; \theta^-) & \text{otherwise} \end{cases} \quad (1)$$

- vi. Perform a gradient descent step on the loss:

$$L(\theta) = (y_j - Q(s_j, a_j; \theta))^2 \quad (2)$$

- vii. Every  $C$  steps, update target network:  $\theta^- \leftarrow \theta$ .

## 4.2 Advantages

- **Efficiency:** By reusing past experiences, the agent makes better use of the data it has collected, leading to improved sample efficiency.
- **Stability:** Random sampling of experiences breaks the correlation between consecutive experiences, leading to more stable and robust learning.

## 4.3 Disadvantages

- **Memory Usage:** Storing a large number of experiences requires significant memory.
- **Complexity:** Implementing and managing the replay buffer adds complexity to the algorithm.

# 5 Training Network

The Q-network is a neural network that approximates the Q-function. It takes the current state as input and outputs the Q-values for all possible actions.

## 5.1 Training Process

1. **Initialization:** Initialize the Q-network with random weights.
2. **Experience Collection:** Interact with the environment to collect experiences  $(s, a, r, s')$ .
3. **Sampling:** Sample a random mini-batch of experiences from the replay buffer.

4. **Target Calculation:** For each experience in the mini-batch, calculate the target Q-value using the Bellman equation.
5. **Loss Calculation:** Calculate the loss between the predicted Q-values and the target Q-values.
6. **Optimization:** Perform a gradient descent step to minimize the loss and update the Q-network weights.

## 5.2 Advantages

- **Generalization:** The neural network can generalize from past experiences to unseen states.
- **Scalability:** DQN can handle high-dimensional state spaces, such as images, through the use of deep convolutional networks.

## 5.3 Disadvantages

- **Stability:** Training deep neural networks can be unstable and may require careful tuning of hyperparameters.
- **Sample Inefficiency:** DQN requires a large number of training samples to learn effectively.

# 6 Conclusion

DQN is a powerful reinforcement learning algorithm that combines the strengths of Q-learning and deep learning. By using a replay buffer and a training network, DQN achieves stable and efficient learning in complex environments. However, it also introduces challenges such as memory usage, hyperparameter tuning, and sample inefficiency.